


# Project3

March 20, 2020

```
[1]: import cv2
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
from utils import *
import os
from scipy.sparse import lil_matrix
from scipy.sparse import linalg
from scipy.sparse import csr_matrix
import pdb
import time
```

```
[2]: img = cv2.imread('samples/toy_problem.png')
toy_img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
toy_img = cv2.cvtColor(toy_img, cv2.COLOR_BGR2GRAY).astype('double') / 255.0
plt.imshow(toy_img, cmap="gray")
```

```
[2]: <matplotlib.image.AxesImage at 0x7fc0588fcd68>
```



## 0.1 Part 1 Toy Problem (20 pts)

```
[3]: def toy_reconstruct(toy_img):  
    """  
    The implementation for gradient domain processing is not complicated, but  
    → it is easy to make a mistake, so let's start with a toy example. Reconstruct  
    → this image from its gradient values, plus one pixel intensity. Denote the  
    → intensity of the source image at (x, y) as s(x,y) and the value to solve for  
    → as v(x,y). For each pixel, then, we have two objectives:  
    1. minimize (v(x+1,y)-v(x,y) - (s(x+1,y)-s(x,y)))2  
    2. minimize (v(x,y+1)-v(x,y) - (s(x,y+1)-s(x,y)))2  
    Note that these could be solved while adding any constant value to v, so we  
    → will add one more objective:  
    3. minimize (v(1,1)-s(1,1))2  
  
    :param toy_img: numpy.ndarray  
    """  
    start = time.time()  
    im_h, im_w = toy_img.shape  
    im2var = np.arange(im_h * im_w).reshape(im_h, im_w) # value is transposed??  
  
    # print("im_h, im_w ", im_h, im_w)  
    # print("im2var shape ", im2var.shape)  
    # print("toy_img shape ", toy_img.shape)  
    # Objective 1  
    # e = e + 1  
    # A[e][im2var[y][x+1]] = 1  
    # A[e][im2var[y][x]] = -1  
    # b[e] = im[y][x+1] - im[y][x]  
  
    # A matrix bounds im_h * im_w  
    # https://piazza.com/class/k5cumohrew35en?cid=362 ??  
    n_constraints = 2 * im_h * im_w + 1 #number of constraints  
    n_pixels = im_h * im_w  
    A = lil_matrix((n_constraints, n_pixels), dtype = np.float64)  
    # print("A ", A.shape)  
    b = np.zeros(n_constraints, np.float64)  
    # print("b ", b.shape)  
  
    e = 0  
    A[e, im2var[0][0]] = 1  
    b[e] = toy_img[0][0]  
    # pdb.set_trace()  
    for y in range(im_h): #im2var.shape[0])
```

```

for x in range(im_w): #im2var.shape[1])

    # Need to handle border cases. x == im_w-1 and y == im_h-1
    #Objective 1
    e = e + 1
    if(x != im_w-1):
        A[e, im2var[y][x+1]] = 1
        A[e, im2var[y][x]] = -1
        b[e] = toy_img[y][x+1] - toy_img[y][x]
    else:
        A[e, im2var[y][x]] = -1
        b[e] = -toy_img[y][x]

    #Objective 2
    e = e + 1
    if(y != im_h-1):
        A[e, im2var[y+1][x]] = 1
        A[e, im2var[y][x]] = -1
        b[e] = toy_img[y+1][x] - toy_img[y][x]
    else:
        A[e, im2var[y][x]] = -1
        b[e] = -toy_img[y][x]

v = linalg.lsqr(csr_matrix(A), b)
res = v[0].reshape(im_h, im_w)
print("Total time taken ", time.time() - start)
return res

```

```

[4]: im_out = toy_reconstruct(toy_img)
    # print("toy_img ", toy_img)
    # print("im_out ", im_out)
    if im_out.any():
        print("Error is: ", np.sqrt(((im_out - toy_img)**2).sum()))

```


Total time taken 1.4244844913482666  
Error is: 0.0001901762385198699

```

[5]: plt.imshow(im_out, cmap='gray')

[5]: <matplotlib.image.AxesImage at 0x7fc058098390>

```



## 0.2 Preparation

```
[68]: # Feel free to change image
# background_img = cv2.cvtColor(cv2.imread('samples/im2.jpg'), cv2.
#   →COLOR_BGR2RGB).astype('double') / 255.0
# background_img = cv2.cvtColor(cv2.imread('samples/tajmahal.jpg'), cv2.
#   →COLOR_BGR2RGB).astype('double') / 255.0
# background_img = cv2.cvtColor(cv2.imread('samples/vtz.jpg'), cv2.
#   →COLOR_BGR2RGB).astype('double') / 255.0
background_img = cv2.cvtColor(cv2.imread('samples/wall.jpg'), cv2.
#   →COLOR_BGR2RGB).astype('double') / 255.0
plt.figure()
plt.imshow(background_img)
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

[68]: <matplotlib.image.AxesImage at 0x7fc0581bd9b0>

```
[70]: # Feel free to change image
# object_img = cv2.cvtColor(cv2.imread('samples/penguin-chick.jpeg'), cv2.
#   →COLOR_BGR2RGB).astype('double') / 255.0
```

```

object_img = cv2.cvtColor(cv2.imread('samples/dog2.jpg'), cv2.COLOR_BGR2RGB).
    ↳astype('double') / 255.0
import matplotlib.pyplot as plt
%matplotlib notebook
mask_coords = specify_mask(object_img)

```

If it doesn't get you to the drawing mode, then rerun this function again.

<IPython.core.display.Javascript object>


<IPython.core.display.HTML object>

```

[71]: xs = mask_coords[0]
      ys = mask_coords[1]
      %matplotlib inline
      import matplotlib.pyplot as plt
      plt.figure()
      mask = get_mask(ys, xs, object_img)

```

<Figure size 432x288 with 0 Axes>



```

[72]: %matplotlib notebook
      import matplotlib.pyplot as plt

```


```
bottom_center = specify_bottom_center(background_img)
```

If it doesn't get you to the drawing mode, then rerun this function again. Also, make sure the object fill fit into the background image. Otherwise it will crash

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
[73]: %matplotlib inline
import matplotlib.pyplot as plt
cropped_object, object_mask = align_source(object_img, mask, background_img,
↳bottom_center)
```



### 0.3 Part 2 Poisson Blending (50 pts)

```
[80]: def poisson_blend(cropped_object, object_mask, background_img):
      """
      :param cropped_object: numpy.ndarray One you get from align_source
      :param object_mask: numpy.ndarray One you get from align_source
      :param background_img: numpy.ndarray
      """
      # print("cropped_object: ", cropped_object.shape)
```

```

#     print("object_mask: ", object_mask.shape)
#     print("background_mask: ", background_img.shape)

# We need to do for 3 channels
# find co-ordinates of non-zero values in object_mask

start = time.time()
im = cropped_object.copy()
co_h, co_w = cropped_object.shape[0], cropped_object.shape[1]

# Generate co2var matrix
co2var = np.arange(co_h * co_w).reshape(co_h, co_w)

# number of pixels we need to process from object mask
mask_coords = np.argwhere(object_mask != 0)
n_constraints = 4 * len(mask_coords) + 1
n_pixels = co_h * co_w

# final result
res3Channels = np.zeros((co_h, co_w, 3))
for ch in range(3): # for each channel

    start_time = time.time()
    print("processing channel ", ch)
    # Generata a A matrix n_constraints, n_pixels
    A = lil_matrix((n_constraints, n_pixels), dtype = np.float64)
    # Generate B Matrix with n_constraints
    b = np.zeros(n_constraints, np.float64)

    #https://piazza.com/class/k5cumohrew35en?cid=451
    #https://piazza.com/class/k5cumohrew35en?cid=466
    #https://piazza.com/class/k5cumohrew35en?cid=484
    print("time spent ", time.time() - start_time)
    start_time = time.time()
    print("apply constraints ")
    e = 0
    for (y, x) in mask_coords:

        if(object_mask[y][x+1] == 1): #x+1
            A[e, co2var[y][x]] = 1
            A[e, co2var[y][x+1]] = -1
            b[e] = im[y][x][ch] - im[y][x+1][ch]
        else:
            A[e, co2var[y][x]] != 1
            b[e] = im[y][x][ch] - im[y][x+1][ch] + 1
    background_img[y][x+1][ch]
#         b[e] = background_img[y][x+1][ch]

```

```

    e = e + 1
    if(object_mask[y+1][x] == 1): #y+1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y+1][x]] = -1
        b[e] = im[y][x][ch] - im[y+1][x][ch]
    else:
        A[e, co2var[y][x]] = 1
        b[e] = im[y][x][ch] - im[y+1][x][ch] +
→background_img[y+1][x][ch]
#         b[e] = background_img[y+1][x][ch]

    e = e + 1
    if (object_mask[y][x-1] == 1): #x-1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y][x-1]] = -1
        b[e] = im[y][x][ch] - im[y][x-1][ch]
    else:
        A[e, co2var[y][x]] = 1
        b[e] = im[y][x][ch] - im[y][x-1][ch] +
→background_img[y][x-1][ch]
#         b[e] = background_img[y][x-1][ch]

    e = e + 1
    if (object_mask[y-1][x] == 1): #y-1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y-1][x]] = -1
        b[e] = im[y][x][ch] - im[y-1][x][ch]
    else:
        A[e, co2var[y][x]] = 1
        b[e] = im[y][x][ch] - im[y-1][x][ch] +
→background_img[y-1][x][ch]
#         b[e] = background_img[y-1][x][ch]
    e = e + 1

    print("calc lsqr ", time.time() - start_time)
    start_time = time.time()
    v = linalg.lsqr(csr_matrix(A), b)
    res = v[0].reshape(co_h, co_w)
    res3Channels[:, :, ch] = cv2.add(res, (1-object_mask)*background_img[:, :,
→,ch])
    print("channel ", ch , " processed in ", time.time() - start_time)

    print("poisson_blend total time ", time.time() - start)
    return res3Channels

```




```
[81]: im_blend = poisson_blend(cropped_object, object_mask, background_img)
      if im_blend.any():
          %matplotlib inline
          import matplotlib.pyplot as plt
          plt.imshow(im_blend)
```

```
processing channel 0
time spent 0.1333932876586914
apply constraints
calc lsqr 0.9930562973022461
channel 0 processed in 1.8847289085388184
processing channel 1
time spent 0.04763150215148926
apply constraints
calc lsqr 0.956810474395752
channel 1 processed in 1.779193639755249
processing channel 2
time spent 0.12368059158325195
apply constraints
calc lsqr 0.9884536266326904
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
channel 2 processed in 1.8233060836791992
poisson_blend total time 5.802590370178223
```



## 0.4 Part 3 Mixed Gradients (20 pts)

```
[82]: def dij(si, sj, ti, tj):
#b[e] = im[y][x][ch] - im[y][x-1][ch] + background_img[y][x-1][ch]
#b[e] = im[y][x][ch] - im[y][x-1][ch]
    if (abs(si - sj) > abs(ti - tj)):
        return (si - sj)
    else:
        return (ti - tj)
#    return si - sj

def mix_blend(cropped_object, object_mask, background_img):
    """
    :param cropped_object: numpy.ndarray One you get from align_source
    :param object_mask: numpy.ndarray One you get from align_source
    :param background_img: numpy.ndarray
    """
    #    print("cropped_object: ", cropped_object.shape)
    #    print("object_mask: ", object_mask.shape)
    #    print("background_mask: ", background_img.shape)

    # We need to do for 3 channels
    #find co-ordinates of non-zero values in object_mask

    start = time.time()
    im = cropped_object.copy()
    co_h, co_w = cropped_object.shape[0], cropped_object.shape[1]

    # Generate co2var matrix
    co2var = np.arange(co_h * co_w).reshape(co_h, co_w)

    # number of pixels we need to process from object mask
    mask_coords = np.argwhere(object_mask != 0)
    n_constraints = 4 * len(mask_coords) + 1
    n_pixels = co_h * co_w

    # final result
    res3Channels = np.zeros((co_h, co_w, 3))
    for ch in range(3): # for each channel

        start_time = time.time()
        print("processing channel ", ch)
        # Generata a A matrix n_constraints, n_pixels
        A = lil_matrix((n_constraints, n_pixels), dtype = np.float64)
        # Generate B Matrix with n_constraints
        b = np.zeros(n_constraints, np.float64)
```

```

#https://piazza.com/class/k5cumohrew35en?cid=451 almost code;
#https://piazza.com/class/k5cumohrew35en?cid=466
#https://piazza.com/class/k5cumohrew35en?cid=484
print("time spent ", time.time() - start_time)
start_time = time.time()
print("apply constraints ")
e = 0
for (y, x) in mask_coords:

    bb = dij(im[y][x][ch], im[y][x+1][ch], background_img[y][x][ch],
→background_img[y][x+1][ch])
    if(object_mask[y][x+1]) > 0: #x+1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y][x+1]] = -1
        b[e] = im[y][x][ch] - im[y][x+1][ch]
        b[e] = bb
    else:
        A[e, co2var[y][x]] = 1
        b[e] = bb + background_img[y][x+1][ch]
    e = e + 1

    bb = dij(im[y][x][ch], im[y+1][x][ch], background_img[y][x][ch],
→background_img[y+1][x][ch])
    if(object_mask[y+1][x]) > 0: #y+1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y+1][x]] = -1
        b[e] = bb
    else:
        A[e, co2var[y][x]] = 1
        b[e] = bb + background_img[y+1][x][ch]
    e = e + 1

    bb = dij(im[y][x][ch], im[y][x-1][ch], background_img[y][x][ch],
→background_img[y][x-1][ch])
    if (object_mask[y][x-1]) > 0: #x-1
        A[e, co2var[y][x]] = 1
        A[e, co2var[y][x-1]] = -1
        b[e] = bb
    else:
        A[e, co2var[y][x]] = 1
        b[e] = bb + background_img[y][x-1][ch]
    e = e + 1

    bb = dij(im[y][x][ch], im[y-1][x][ch], background_img[y][x][ch],
→background_img[y-1][x][ch])
    if (object_mask[y-1][x] > 0): #y-1
        A[e, co2var[y][x]] = 1

```

```

        A[e, co2var[y-1][x]] = -1
        b[e] = bb
    else:
        A[e, co2var[y][x]] = 1
        b[e] = bb + background_img[y-1][x][ch]
    e = e + 1

    print("calc lsqr ", time.time() - start_time)
    start_time = time.time()
    v = linalg.lsqr(csr_matrix(A), b)
    res = v[0].reshape(co_h, co_w)
    res3Channels[:, :, ch] = cv2.add(res, (1-object_mask)*background_img[:, :,
→, ch])
    print("channel ", ch, " processed in ", time.time() - start_time)

    print("fix last mix_blend total time ", time.time() - start)
    return res3Channels

```

```

[83]: im_mix = mix_blend(cropped_object, object_mask, background_img)
    if im_mix.any():
        %matplotlib inline
        import matplotlib.pyplot as plt
        plt.imshow(im_mix)

```

```

processing channel 0
time spent 0.042109012603759766
apply constraints
calc lsqr 1.202021837234497
channel 0 processed in 2.1095380783081055
processing channel 1
time spent 0.1224663257598877
apply constraints
calc lsqr 1.2107105255126953
channel 1 processed in 2.1004998683929443
processing channel 2
time spent 0.04578566551208496
apply constraints
calc lsqr 1.1856167316436768


```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```

channel 2 processed in 2.068063974380493
fix last mix_blend total time 6.495405912399292

```



## 1 Bells & Whistles (Extra Points)

### 1.1 Color2Gray (20 pts)

```
[ ]: color_img = cv2.imread('samples/colorBlind8.png')
plt.imshow(color_img)

[ ]: gray_img = cv2.cvtColor(color_img, cv2.COLOR_BGR2GRAY)
plt.imshow(gray_img, cmap="gray")

[ ]: #taking the max among the 3 color channels solved it for me.
def gradient_norm(si, sj):
    max = 0
    for ch in range(3):
        if abs(si[ch] - sj[ch]) > abs(max):
            max = si[ch] - sj[ch]
    return max

def color2gray(img):
    im_h, im_w, _ = img.shape
    im2var = np.arange(im_h * im_w).reshape(im_h, im_w)

    n_constraints = 2 * im_h * im_w + 1 #number of constraints
    n_pixels = im_h * im_w
    A = lil_matrix((n_constraints, n_pixels), dtype = np.float64)
```

```

b = np.zeros(n_constraints, np.float64)

e = 0
A[e, im2var[0,0]] = 1
b[e] = gray_img[0,0]

for y in range(im_h):
    for x in range(im_w):

        # Need to handle border cases. x == im_w-1 and y == im_h-1
        #Objective 1
        e = e + 1
        if(x != im_w-1):
            A[e, im2var[y][x+1]] = -1
            A[e, im2var[y][x]] = 1
            # rgb values are in here..
            # similar gradients to the original RGB
            # each value has 3 channels
            b[e] = gradient_norm(img[y,x], img[y,x+1])
        else:
            A[e, im2var[y][x]] = 1
#            b[e] = -gray_img[y][x]

#            pdb.set_trace()
        #Objective 2
        e = e + 1
        if(y != im_h-1):
            A[e, im2var[y+1][x]] = -1
            A[e, im2var[y][x]] = 1
            b[e] = gradient_norm(img[y,x], img[y+1,x])
        else:
            A[e, im2var[y][x]] = -1
#            b[e] = -gray_img[y][x]

v = linalg.lsqr(csr_matrix(A), b)
res = v[0].reshape(im_h, im_w)
return res

grayed_image = color2gray(color_img.astype('float32'))
plt.imshow(grayed_image, cmap='gray')

```

## 1.2 Laplacian pyramid blending (20 pts)

```

[85]: def laplacian_blend(img1, img2):

        mask = np.zeros_like(cropped_object)

```

```

# create mask for 3 channels
for i in range(3):
    mask[:, :, i] = object_mask.copy()

levels = 121
# generate a gaussian pyramid for A
gImg1 = img1.copy()
gpImg1 = [gImg1]
for i in range(levels):
    gImg1 = cv2.pyrDown(gImg1)
    gpImg1.append(np.float32(gImg1))

# generate a gaussian pyramid for B
gImg2 = img2.copy()
gpImg2 = [gImg2]
for i in range(levels):
    gImg2 = cv2.pyrDown(gImg2)
    gpImg2.append(np.float32(gImg2))

# generate a gaussian pyramid for Mask
gMask = mask.copy()
gpMask = [gMask]
for i in range(levels):
    gMask = cv2.pyrDown(gMask)
    gpMask.append(np.float32(gMask))
gpMask.reverse()

#generate laplacian pyramid for Img1
lpImg1 = [gpImg1[levels-1]]
for i in range(levels-1, 0, -1):
    G = cv2.pyrUp(gpImg1[i])(:gpImg1[i-1].shape[0], :gpImg1[i-1].shape[1], :)
    L = np.subtract(gpImg1[i-1], G)
    lpImg1.append(L)

#generate laplacian pyramid for Img2
lpImg2 = [gpImg2[levels-1]]
for i in range(levels-1, 0, -1):
    G = cv2.pyrUp(gpImg2[i])(:gpImg2[i-1].shape[0], :gpImg2[i-1].shape[1], :
→]
    L = np.subtract(gpImg2[i-1], G)
    lpImg2.append(L)

alphamatte = []
for a, b, alpha in zip(lpImg1, lpImg2, gpMask[1:]):
    alphamatte.append(a * alpha + b * (1.0 - alpha))

resAllLevels = alphamatte[0]

```



```

for i in range(1, levels):
    lvl = cv2.pyrUp(resAllLevels)
    reslvl = lvl[:alphamatte[i].shape[0],:alphamatte[i].shape[1]]
    resAllLevels = cv2.add(reslvl, np.float32(alphamatte[i]))

return resAllLevels

```

```

[86]: out_lap = laplacian_blend(cropped_object, background_img)
      plt.imshow(out_lap)


```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```

[86]: <matplotlib.image.AxesImage at 0x7fc0582d1fd0>

```



### 1.3 More gradient domain processing (up to 20 pts)